Hybrid Retrieval for Papers

1. Introduction

Retrieval augmented generation grounds a language model in passages found in the document. The hybrid retriever combines BM25 keyword scores with dense embedding similarity [1]. Reciprocal Rank Fusion merges both rankings and reached 94.2% recall in our tests [2, 3]. Chunks overlap by one hundred characters so that sentences cut at a boundary stay readable. Retrieval augmented generation grounds a language model in passages found in the document. The hybrid retriever combines BM25 keyword scores with dense embedding similarity [1]. Reciprocal Rank Fusion merges both rankings and reached 94.2% recall in our tests [2, 3]. Chunks overlap by one hundred characters so that sentences cut at a boundary stay readable.

Second paragraph of this section describes the evaluation with deter-ministic settings and a fixed seed. Retrieval augmented generation grounds a language model in passages found in the document. The hybrid retriever combines BM25 keyword scores with dense embedding similarity [1]. Reciprocal Rank Fusion merges both rankings and reached 94.2% recall in our tests [2, 3]. Chunks overlap by one hundred characters so that sentences cut at a boundary stay readable.

2. Methods

2.1 Chunking

Retrieval augmented generation grounds a language model in passages found in the document. The hybrid retriever combines BM25 keyword scores with dense embedding similarity [1]. Reciprocal Rank Fusion merges both rankings and reached 94.2% recall in our tests [2, 3]. Chunks overlap by one hundred characters so that sentences cut at a boundary stay readable. Retrieval augmented generation grounds a language model in passages found in the document. The hybrid retriever combines BM25 keyword scores with dense embedding similarity [1]. Reciprocal Rank Fusion merges both rankings and reached 94.2% recall in our tests [2, 3]. Chunks overlap by one hundred characters so that sentences cut at a boundary stay readable.

Second paragraph of this section describes the evaluation with deter-ministic settings and a fixed seed. Retrieval augmented generation grounds a language model in passages found in the document. The hybrid retriever combines BM25 keyword scores with dense embedding similarity [1]. Reciprocal Rank Fusion merges both rankings and reached 94.2% recall in our tests [2, 3]. Chunks overlap by one hundred characters so that sentences cut at a boundary stay readable.